

Tutorial 6 Solution

Introduction

For an overview of the basic ideas and background on Bloom filter see its wikipedia entry or the original paper due to Bloom[1]. Here we only enumerate the relevant results required to solve the tutorial problems.

A Bloom filter can be used to answer queries to test the membership of elements of a set X in the form “does $x \in X$ belong to the subset $S \subseteq X$?”. The filter consists of an array of M bits and a collection of K hash functions $f_1, f_2, \dots, f_K : X \mapsto \{0, 1, \dots, M-1\}$. The hash functions are chosen in such a way that (a priori) for each $x \in X$ and each $1 \leq k \leq K$ the probability of mapping x to any of the M bits equals $1/M$.

Initially, all bits in the array are set to zero (which corresponds to $S = \emptyset$). Whenever a new element $x \in X$ is inserted to the filter (corresponding to the operation $S \leftarrow S \cup \{x\}$), one applies all K hash functions to x and sets the bits with indices $f_1(x), f_2(x), \dots, f_K(x)$ to value 1 (true). Note that in the case of a traditional Bloom filter it is not possible to remove elements.

Whenever the filter receives the query “ $x \in S$?” one computes again all values $f_1(x), f_2(x), \dots, f_K(x)$ and reads the value of the corresponding bits in the array. If at least one of these bits (still) has a value 0 (false), the filter returns “ $x \notin S$ ”. Note that this answer is guaranteed to be correct, since if it was the case that x was previously added to S , then all of these bits would have been set to 1 (true). However, in the case that all of them are set to 1, then the filter returns “ $x \in S$ ”, which may be an incorrect answer. If the filter claims that an element is part of S while in fact it is not, then we speak of a *false positive*.

For a given Bloom filter, we are able to obtain the *false-positive probability* based on the assumptions on the size of the array and the number of uniformness properties of the hash functions.

$$\begin{aligned} \text{Prob}(\text{bit } b \text{ is set to 1 by hash function } f_k) &= \frac{1}{M} \\ \text{Prob}(\text{bit } b \text{ is not set to 1 by any hash function } f_k) &= \left(1 - \frac{1}{M}\right)^K \\ \text{Prob}(\text{bit } b \text{ is still set to 0 after inserting } N \text{ elements}) &= \left(1 - \frac{1}{M}\right)^{KN} \end{aligned}$$

Then the false-positive probability can be computed as follows.

$$\begin{aligned} &\text{Prob}(\text{falsely classifying } x \in S \text{ after inserting } N \text{ elements}) \\ &= \text{Prob}(\text{all of the } K \text{ bits } f_1(x), f_2(x), \dots, f_K(x) \text{ are set to 1 after inserting } N \text{ elements}) \\ &= (1 - \text{Prob}(\text{bit } b = f_k(x) \text{ is still set to 0 after inserting } N \text{ elements}))^K \\ &= \left(1 - \left(1 - \frac{1}{M}\right)^{KN}\right)^K \approx (1 - e^{-KN/M})^K \end{aligned} \tag{1}$$

For a fixed array size M and a fixed number of insertions N , one can ask for the (smallest) value of K that minimizes the false-positive probability. Here we only mention the closed form expression for the optimal K , skip the details and instead refer to the literature. So the smallest value of K (denote K^*) that minimizes the false positive probability can be computed as

$$K^* = \frac{M}{N} \ln 2 \approx 0.7 \frac{M}{N}. \tag{2}$$

Assuming one chooses the number of hash functions according to K^* , one obtains from Equations (1) and (2) a false positive probability of approximately

$$(1 - e^{-K^*N/M})^{K^*} = 2^{-\frac{M}{N} \ln 2}. \quad (3)$$

Solutions

1. a) We have $K = 2$, $M = 10^6 \cdot 8$ bits, $N = 10^5$. From (1) we obtain

$$\text{Prob}(\text{falsely classifying } x \in S \text{ after inserting } N \text{ elements}) \approx (1 - e^{-1/40})^2 \approx 0.00061.$$

- b) Now instead we have $N = 10^6$ and thus obtain from (1)

$$\text{Prob}(\text{falsely classifying } x \in S \text{ after inserting } N \text{ elements}) \approx (1 - e^{-1/4})^2 \approx 0.04893.$$

2. a) We have $M = 10^6 \cdot 8$ bits and $N = 10^6$. From (2) we obtain¹

$$K^* \approx 5.6$$

If we substitute for M and N we obtain from (3)

$$\begin{aligned} & \text{Prob}(\text{falsely classifying } x \in S \text{ after inserting } N \text{ elements}) \\ = & 2^{-M/N \cdot \ln 2} \approx 0.021. \end{aligned}$$

- b) Instead we now have $M = 512 \cdot 10^3 \cdot 8$ bits. As before, we use (2) to obtain

$$K^* \approx 2.9$$

and for the false-positive probability

$$\begin{aligned} & \text{Prob}(\text{falsely classifying } x \in S \text{ after inserting } N \text{ elements}) \\ = & 2^{-M/N \cdot \ln 2} \approx 0.140. \end{aligned}$$

References

- [1] Burton H. Bloom. Space/time trade-offs in hash coding with allowable errors. *Commun. ACM*, 13:422–426, July 1970.

¹Of course in reality K^* would need to be an integer but let's ignore this here.